


# Analysis of NGS data

*Principles of genotype and SNP calling  
and estimation of allele frequencies*

Matteo Fumagalli

---

We are  
bioinformaticians  
thats what we do





## Presentation outline


- 1 Introduction
  - 2 Genotype likelihoods
  - 3 Genotype calling
  - 4 SNP calling
  - 5 Imputation
-

# Sanger sequencing

aka first/former generation sequencing



# Next Generation Sequencing



- A. Extracted gDNA  
B. gDNA is fragmented into a library of small segments that are each sequenced in parallel.  
C. Individual sequence reads are reassembled by aligning to a reference genome  
D. The whole-genome sequence is derived from the consensus of aligned reads.

>ARPM2ref|NC\_000001.10|:2938046-2939467 Homo sapiens chromosome 1, GRCh37 primary reference assembly

CCAATGATTTTTTTCCGTGTTTCAGAATACGGTTAA  
+SRR038845.53 HWI-EAS038:6:1:0:1474 length=36  
BCCBA@BB@BBBBAB@B9B@-BABA@A:@693:GB=  
@SRR038845.53 HWI-EAS038:6:1:1:360 length=36  
GTTCAAAAAGAACATAAATGTGTCAATAGAAAACCT  
+SRR038845.53 HWI-EAS038:6:1:1:360 length=36

[illegible]

```
#fileformat=VCFv4.1
#fileDate=20140930
#source=23andme2vcf.pl https://github.com/arrantrobot/23andme2vcf
#reference=HG19/hg19.ref.txt.gz
##FORMAT=ID=GT,Number=1,Type=String,Description="Genotype"
##CHROM POS ID REF ALT QUAL FILTER INFO FORMAT GENOTYP
chr1 82154 rs4477212 . . . . . GT
chr1 752566 rs3894315 G A . . . . GT
/r1
chr1 757271 rs3131972 A G . . . . GT
/r1
chr1 798959 rs1124077 . . . . . GT
/0
chr1 800087 rs6681849 T C . . . . GT
/
```


# Forget about

The image is a composite of three elements:

- Top Left:** The Windows 10 logo in white text on a blue background.
- Top Right:** A screenshot of a Windows 10 desktop environment. The taskbar shows the Start button, a search bar, and several pinned apps including Microsoft Edge, File Explorer, and the Windows Store. The Start menu is open, displaying a list of apps and a search bar. The background is a blue wall with a large, stylized '10'.
- Bottom:** A screenshot of a Microsoft Excel spreadsheet. The spreadsheet contains data for a fleet management system, including driver names, run numbers, run dates, run miles, run gross, driver pay, fuel sur-charge, and fuel miles. The data is organized into columns and rows, with some cells containing formulas and totals.

# Statistical inference


Data (D)



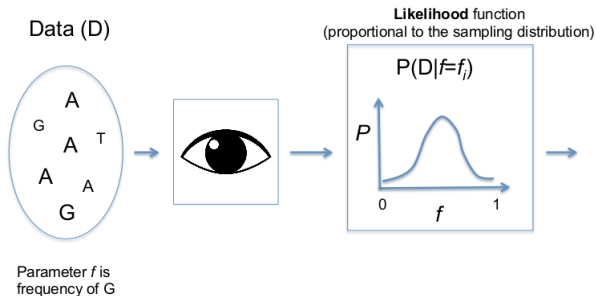
Parameter  $f$  is  
frequency of G



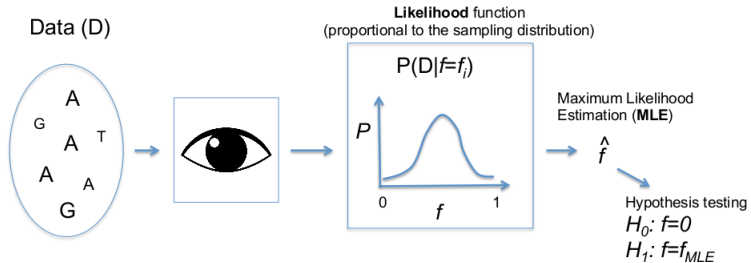
# Statistical inference



# Statistical inference




# Statistical inference



## Likelihood approach:

- All the information on the parameter is in the likelihood function (we use all the data!).
- More data leads to less bias and less variance.
- Suitable for hypothesis testing.

## The data



- is a **nucleotide**/base/allele with a certain **quality** score

## Genotype likelihoods

### Likelihood

$$P(D|G = \{A_1, A_2, \dots, A_n\})$$

with


$A_i \in \{A, C, G, T\}$  and  $n$  being the ploidy

How many genotypes likelihoods do we need to calculate for each individual at each site?

# Genotype likelihoods

Chrom1    272    A    24    AAAAAGGAGAGGTAAG    <<<+;<<<<<<<<<=<;<;7<&

Base quality in Phred scale



## Calculating genotype likelihoods

### Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

- $L_{A_j,i} = P(D|A_G = A_j)$
- $A_i \in \{A, C, G, T\}$
- $R$  is the depth (nr. of reads)
- $N$  is the ploidy (nr. of chromosomes)

Example:

AAAG with all quality scores equal to 20 (in phred score)

$P(D|G = AC) = ?$

## Calculating genotype likelihoods

### Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

AAAG & Q=20

$$P(D|G = \{A, C\}) = \dots$$



## Calculating genotype likelihoods

### Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

AAAG & Q=20

$N = 2; i = 1; A_1 = A; A_2 = C$

$$P(D|G = \{A, C\}) = \left( \frac{L_{A,1}}{2} + \frac{L_{C,1}}{2} \right) \times \dots$$

What are  $L_{A,1}$  and  $L_{C,1}$ ?

## Calculating genotype likelihoods

### Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

AAAG & Q=20

$$L_{C,1} = \frac{\epsilon}{3}$$

$$L_{A,1} = 1 - \epsilon$$

$$P(D|G = \{A, C\}) = \left(\frac{1-\epsilon}{2} + \frac{\epsilon}{6}\right) \times \dots$$

## Calculating genotype likelihoods

### Likelihood function

$$P(D|G = \{A_1, A_2, \dots, A_N\}) = \prod_{i=1}^R \sum_{j=1}^N \frac{L_{A_j,i}}{N}$$

AAAG & Q=20

$$L_{C,1} = \frac{\epsilon}{3}$$

$$L_{A,1} = 1 - \epsilon$$

$$P(D|G = \{A, C\}) = \left(\frac{1-\epsilon}{2} + \frac{\epsilon}{6}\right) \times \frac{\epsilon}{3} \times \frac{\epsilon}{3} \times \frac{\epsilon}{3}$$

What is  $\epsilon$ ?

## Calculating genotype likelihoods

Genotype	Likelihood (log10)
AA	-2.49
<b>AC</b>	<b>-3.38</b>
AG	-1.22
AT	-3.38
CC	-9.91
CG	-7.74
CT	-9.91
GG	-7.44
GT	-7.74
TT	-9.91

AAAG &  $\epsilon = 0.01$

## Genotype calling

Genotype	Likelihood (log10)
AA	-2.49
AC	-3.38
AG	-1.22
AT	-3.38
CC	-9.91
CG	-7.74
CT	-9.91
GG	-7.44
GT	-7.74
TT	-9.91

AAAG &  $\epsilon = 0.01$

What is the genotype here?

## Genotype calling

Genotype	Likelihood (log10)
AA	-2.49
AC	-3.38
<b>AG</b>	<b>-1.22</b>
AT	-3.38
CC	-9.91
CG	-7.74
CT	-9.91
GG	-7.44
GT	-7.74
TT	-9.91

AAAG &  $\epsilon = 0.01$

What is the genotype?

AG.

### Maximum Likelihood

The simplest genotype caller: choose the genotype with the highest likelihood.

## Major and minor alleles

### Likelihood function

$$\log P(D|G = A) = \sum_{i=1}^R \log L_{A_j,i}$$

AAAG &  $\epsilon = 0.01$

Allele	Likelihood
<b>A</b>	<b>-2.49</b>
C	-3.38
<b>G</b>	<b>-1.22</b>
T	-3.38

We can reduce the genotype space to 3 entries (from 10).

## Genotype calling

AAAG &  $\epsilon = 0.01$  & A,G alleles

Genotype	Likelihood
AA	-5.73
AG	-2.80
GG	-17.12

Examples varying qualities and reads...



## Genotype likelihood ratio

$$\log_{10} \frac{L_{G(1)}}{L_{G(2)}} > t$$

i.e.  $t = 1$  meaning that the most likely genotype is 10 times more likely than the second most likely one

Pros and cons?

- Yes:
- No:

## Genotype likelihood ratio

$$\log_{10} \frac{L_{G(1)}}{L_{G(2)}} > t$$

i.e.  $t = 1$  meaning that the most likely genotype is 10 times more likely than the second most likely one

Pros and cons?

- Yes: genotype are called with higher **confidence**
- No: more **missing** data

# Statistical thinking



Figure 1: Nessie, the Loch Ness Monster. True or fake?

## Statistical thinking

- $T = \{0, 1\}$ , whether I tell you I saw Nessie or not.
- $N = \{0, 1\}$ , whether Nessie exists (I saw it) or not.

### Questions

- What are  $p(T = 1|N = 1)$  and  $p(T = 1|N = 0)$ ?
- What is a Maximum Likelihood Estimate of  $N$ ?

## Statistical thinking

Our inference on  $N$  is driven solely by our observations, given by our likelihood function.




Figure 2: The eye: a "likelihood" organ.

## Statistical thinking

In real life we take many decisions based not only on what we observe but also on some believes of ours.




Figure 3: The brain: a "non-likelihood" organ.

## Bayesian thinking

- with "eyes only" our intuition is that  $p(N|T) \approx p(T|N)$
- with "the brain" our intuition is that  $p(N|T) \approx p(T|N)p(N)$

Our "belief" expresses the probability  $p(N)$  **unconditional** of the data.

### Question

How can we define  $p(N)$ ?

## Bayesian thinking

The "belief" function  $p(N)$  is called **prior probability** and the joint product of the likelihood  $p(T|N)$  and the prior is proportional to the **posterior probability**  $p(N|T)$ .

The use of posterior probabilities for inferences is called Bayesian statistics.



## Statistical inference

If  $D$  is the data and  $\theta$  is your unknown parameter, then

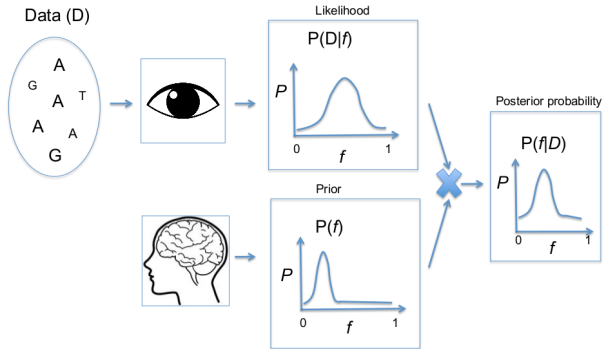
- the frequentist conditions on parameters and integrates over the data,  $p(D|\theta)$ ,
- the Bayesian conditions on the data and integrates over the parameters,  $p(\theta|D)$ .

# Statistical inference

## Bayesian vs. Likelihoodist

- we derive "proper" probability distributions of our parameters rather than deriving a point estimate;
- a probability is assigned to a hypothesis rather than a hypothesis is tested;
- we can "accept" the null hypothesis rather than "fail to reject" it;
- parsimony imposed in model choice rather than correcting for multiple tests.

# Bayesian inference



## Bayesian concepts

### Bayes' Theorem

$$p(\vec{\theta}|\vec{y}) = \frac{f(\vec{y}|\vec{\theta})\pi(\vec{\theta})}{m(\vec{y})} = \frac{f(\vec{y}|\vec{\theta})\pi(\vec{\theta})}{\int f(\vec{y}|\vec{\theta})d\vec{\theta}} \quad (1)$$

- $\vec{\theta}$  is not a fixed parameter but a random quantity with prior distribution  $\pi(\vec{\theta})$
- $p(\vec{\theta}|\vec{y})$  is the posterior probability distribution of  $\vec{\theta}$
- $\int p(\vec{\theta}|\vec{y})d\vec{\theta} = 1$

## Genotype posterior probability

AAAG &  $\epsilon = 0.01$  & A,G alleles

Genotype	Likelihood (log)	Prior	Posterior
AA	-5.73		
AG	-2.80		
GG	-17.12		

## Genotype posterior probability

AAAG &  $\epsilon = 0.01$  & A,G alleles

Genotype	Likelihood (log)	Prior	Posterior
AA	-5.73	1/3	0.05
AG	-2.80	1/3	0.95
GG	-17.12	1/3	0

Only call genotypes if the largest probability is above a certain threshold (e.g. 0.95).

## Genotype posterior probability

AAAG &  $\epsilon = 0.01$  & A,G alleles & **A is the reference allele**  
 $P(AA) > P(AG) > P(GG)$

Genotype	Likelihood (log)	Prior	Posterior
AA	-5.73	0.80	0.22
AG	-2.80	0.15	0.78
GG	-17.12	0.05	0

The reference allele is just one of the possible alleles, often chosen arbitrarily: why give it so much weight?

## Genotype posterior probability

AAAG &  $\epsilon = 0.01$  & A,G alleles &  $f(A) = 0.7$  from a reference panel

$P(AA) = ?$ ;  $P(AG) = ?$ ;  $P(GG) = ?$

Genotype	Likelihood (log)	Prior	Posterior
AA	-5.73		
AG	-2.80		
GG	-17.12		



## Genotype posterior probability


AAAG &  $\epsilon = 0.01$  & A,G alleles &  $f(A) = 0.7$  from a reference panel

$P(AA) = ?$ ;  $P(AG) = ?$ ;  $P(GG) = ?$

Genotype	Likelihood (log)	Prior	Posterior
AA	-5.73	0.49	0.06
AG	-2.80	0.42	0.94
GG	-17.12	0.09	0

If the assumption of HWE can be reasonably met.

# Empirical Bayesian inference



## Genotype posterior probability

AAAG &  $\epsilon = 0.01$  & A,G alleles &  $f(A) = 0.6$  from the data itself  
 $P(AA) = ?$ ;  $P(AG) = ?$ ;  $P(GG) = ?$


Genotype	Likelihood (log)	Prior	Posterior
AA	-5.73	0.49	0.04
AG	-2.80	0.42	0.96
GG	-17.12	0.09	0

- if the assumption of HWE can be reasonably met
- if you have enough samples to have a robust estimate of the allele frequencies

How can we estimate allele frequencies?

# SNP calling procedures

- Alignment-based caller



We completely rely on how reads have been mapped

Figure from Erik Garrison

## SNP calling procedures

- Assembly-based caller (as in GATK)

Local re-alignment around putative variants; better resolution for INDELs detection.

- Haplotype-based caller (as in freebayes)




Figure from Erik Garrison

## Estimating allele frequencies

Assuming 2 alleles (A,G) with true allele frequency of 0.50

Sample	True genotype	Reads allele A	Read allele G
1	AA	7	0
2	AA	25	1
3	AG	5	3
4	AG	4	4
5	GG	0	2
6	GG	0	4

What is the simplest estimator of allele frequencies?

## Estimating allele frequencies

Assuming 2 alleles (A,G) with true allele frequency of 0.50

Sample	True genotype	Reads allele A	Read allele G
1	AA	7	0
2	AA	25	1
3	AG	5	3
4	AG	4	4
5	GG	0	2
6	GG	0	4
Total		41	14

$$\hat{f} = \frac{\sum_{i=1}^N n_{A,i}}{\sum_{i=1}^N (n_{A,i} + n_{G,i})}$$

$$\hat{f} = 0.75$$

What is wrong with this estimator?

## Estimating allele frequencies

Assuming 2 alleles (A,G) with true allele frequency of 0.50

Sample	True genotype	Reads allele A	Read allele G
1	AA	7	0
2	AA	25	1
3	AG	5	3
4	AG	4	4
5	GG	0	2
6	GG	0	4
Total		41	14

$$\hat{n}_A = \sum_{i=1}^N (1 - \epsilon)n_{A,i} + \epsilon n_{G,i} - \epsilon n_{A,i} - (1 - \epsilon)n_{G,i}$$

$$\hat{f} = 0.77$$



## Estimating allele frequencies

Maximum Likelihood estimator

$$P(D|f) = \prod_{i=1}^N \sum_{g \in \{0,1,2\}} P(D|G = g)P(G = g|f)$$

## Estimating allele frequencies

### Maximum Likelihood estimator

$$P(D|f) = \prod_{i=1}^N \sum_{g \in \{0,1,2\}} P(D|G = g)P(G = g|f)$$

$P(D|G = g)$  is the genotype likelihood and  $P(G = g|f)$  is given by HWE (for instance).

In our previous example,  $\hat{f} = 0.46$  which is much closer to the true value than previous estimators.

# SNP calling

## Challenges

- If high levels of missing data, then genotypes can be lost.
- Rare variants are hard to detect.
- Trade off between false positive and false negative rates.

## How to call SNPs?

- If at least one heterozygous genotype has been called.
- If the estimated allele frequency is above a certain threshold.

## SNP calling

Call a SNP if

$$\hat{f} \geq t$$

where  $t$  can be the minimum sample allele frequency detectable (e.g.  $t = 1/2N$  with  $N$  diploids).

## Likelihood Ratio Test

A Likelihood Ratio Test (LRT) compares the goodness of fit between the null and the alternative model:


- Null model:  $f = 0$
- Alternative model:  $f \neq 0$

$$T = -2 \log \frac{L(f = 0)}{L(f = \hat{f}_{MLE})}$$

where  $T$  is  $\chi^2$  distributed with 1 degree of freedom.

# Haplotype imputation

## Haplotype imputation - simplified



### Reference

- 1000 Genomes
- Phased using family structures

### new data

- partial information

### Imputed data

- Probabilistic approach
- The results retains the uncertainty of both the genotype and the haplotypes

# Haplotype imputation

## Haplotype imputation - simplified

Reference data



New data



Imputed data



### Reference

- 1000 Genomes
- Phased using family structures

### new data

- Data with known and unknown genotypes

### Imputed data

$$p(? = T) =$$

$$p(? = A) =$$

# Haplotype imputation

## Haplotype imputation - simplified



### Reference

- haplotype frequencies

### new data

- Data with known and unknown genotypes

### first haplotype

$$p(? = T) = \frac{0.56}{0.56+0.03} = 0.95$$

$$p(? = A) = \frac{0.03}{0.56+0.03} = 0.05$$

### second haplotype

$$p(? = T) = \frac{0.21}{0.21+0.2} = 0.51$$

$$p(? = A) = \frac{0.2}{0.21+0.2} = 0.49$$

Anders Albrechtsen



# Haplotype imputation

## Haplotype imputation - simplified



### Bayes formula

$$p(H = h|f, G) = \frac{P(G|H=h)P(H=h|f)}{\sum_{h'} P(G|H=h')P(H=h'|f)}$$

### $P(G|H = h)$


1 if consistent

0 otherwise

### first haplotype

$$p(? = T) = \frac{0.56}{0.56+0.03} = 0.95$$

$$p(? = A) = \frac{0.03}{0.56+0.03} = 0.05$$



Raghavan et al. 2015 Science

Thank you for your attention

---